

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
sns.set_theme(color_codes=True)
pd.set_option('display.max_columns', None)
```

```
In [2]: df = pd.read_csv('airline.csv')
df.head()
```

Out[2]:

	Unnamed: 0	id	Gender	Customer Type	Age	Type of Travel	Class	Flight Distance	Inflight wifi service	Departure/ time conv
0	0	70172	Male	Loyal Customer	13	Personal Travel	Eco Plus	460	3	
1	1	5047	Male	disloyal Customer	25	Business travel	Business	235	3	
2	2	110028	Female	Loyal Customer	26	Business travel	Business	1142	2	
3	3	24026	Female	Loyal Customer	25	Business travel	Business	562	2	
4	4	119299	Male	Loyal Customer	61	Business travel	Business	214	3	

Remove ID Column

```
In [3]: df.drop(columns=['Unnamed: 0', 'id'], inplace=True)
df.shape
```

Out[3]: (103904, 23)

Exploratory Data Analysis

```

In [4]: # List of categorical variables to plot
cat_vars = ['Gender', 'Customer Type', 'Type of Travel', 'Class']


# create figure with subplots
fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(15, 15))
axs = axs.flatten()

# create barplot for each categorical variable
for i, var in enumerate(cat_vars):
    sns.countplot(x=var, hue='satisfaction', data=df, ax=axs[i])
    axs[i].set_xticklabels(axs[i].get_xticklabels(), rotation=90)

# adjust spacing between subplots
fig.tight_layout()

# show plot
plt.show()

```



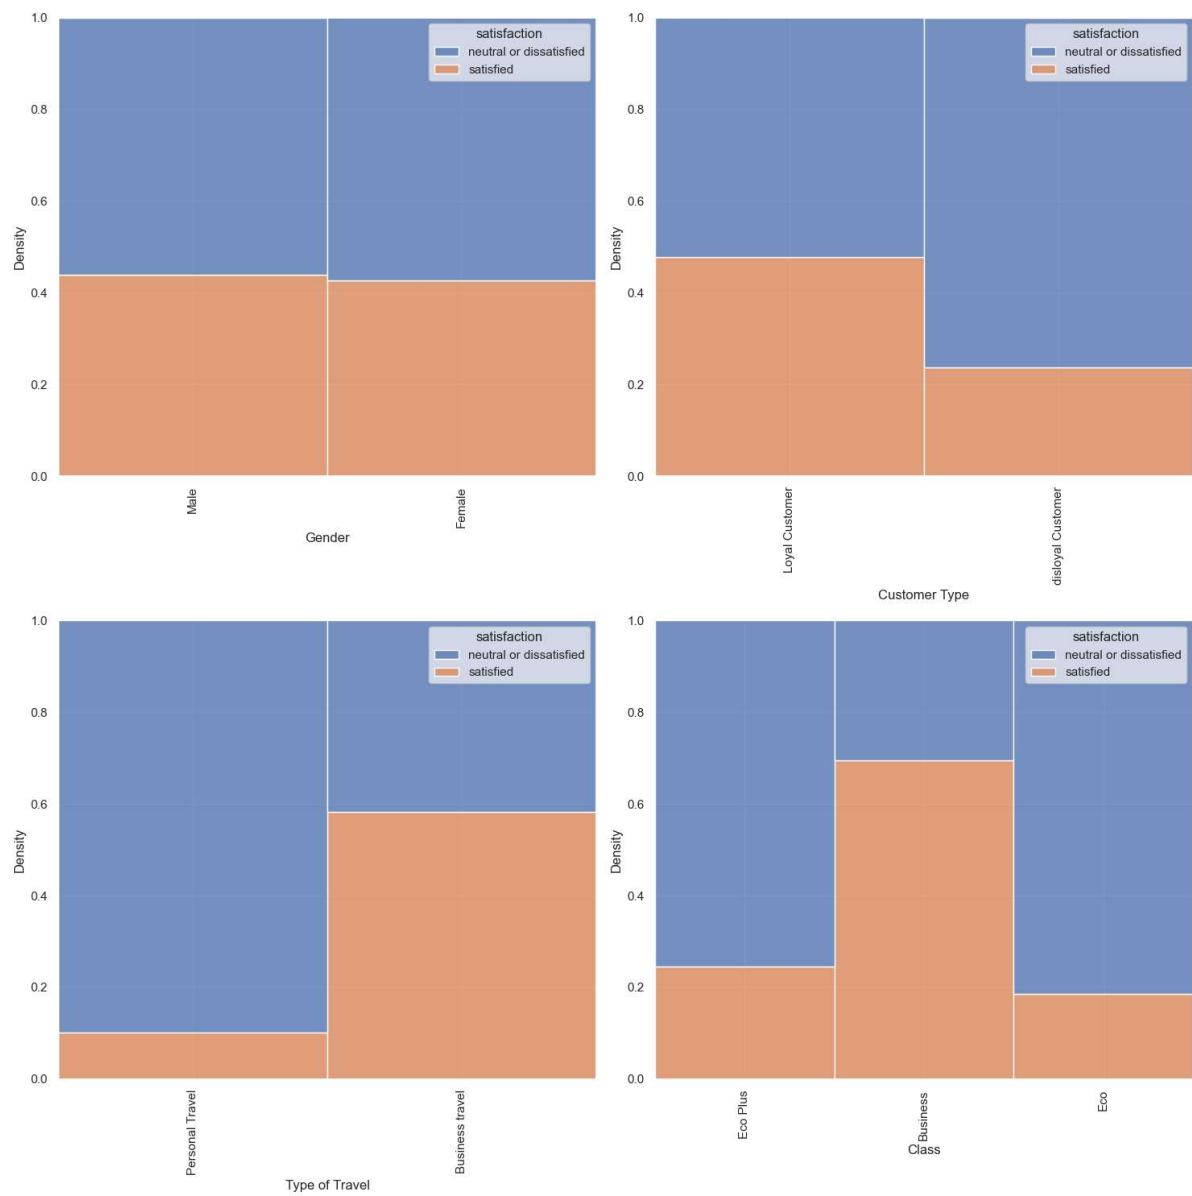
```
In [5]: import warnings
warnings.filterwarnings("ignore")
# get list of categorical variables
cat_vars = ['Gender', 'Customer Type', 'Type of Travel', 'Class']

# create figure with subplots
fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(15, 15))
axs = axs.flatten()

# create histplot for each categorical variable
for i, var in enumerate(cat_vars):
    sns.histplot(x=var, hue='satisfaction', data=df, ax=axs[i], multiple="fill")
    axs[i].set_xticklabels(df[var].unique(), rotation=90)
    axs[i].set_xlabel(var)

# adjust spacing between subplots
fig.tight_layout()

# show plot
plt.show()
```



```

In [6]: cat_vars = ['Inflight wifi service', 'Departure/Arrival time convenient', 'Ease
                  'Gate location', 'Food and drink', 'Online boarding', 'Seat comfort',
                  'On-board service', 'Leg room service', 'Baggage handling', 'Checkin
                  'Cleanliness']

# create a figure and axes
fig, axs = plt.subplots(nrows=3, ncols=5, figsize=(15, 15))

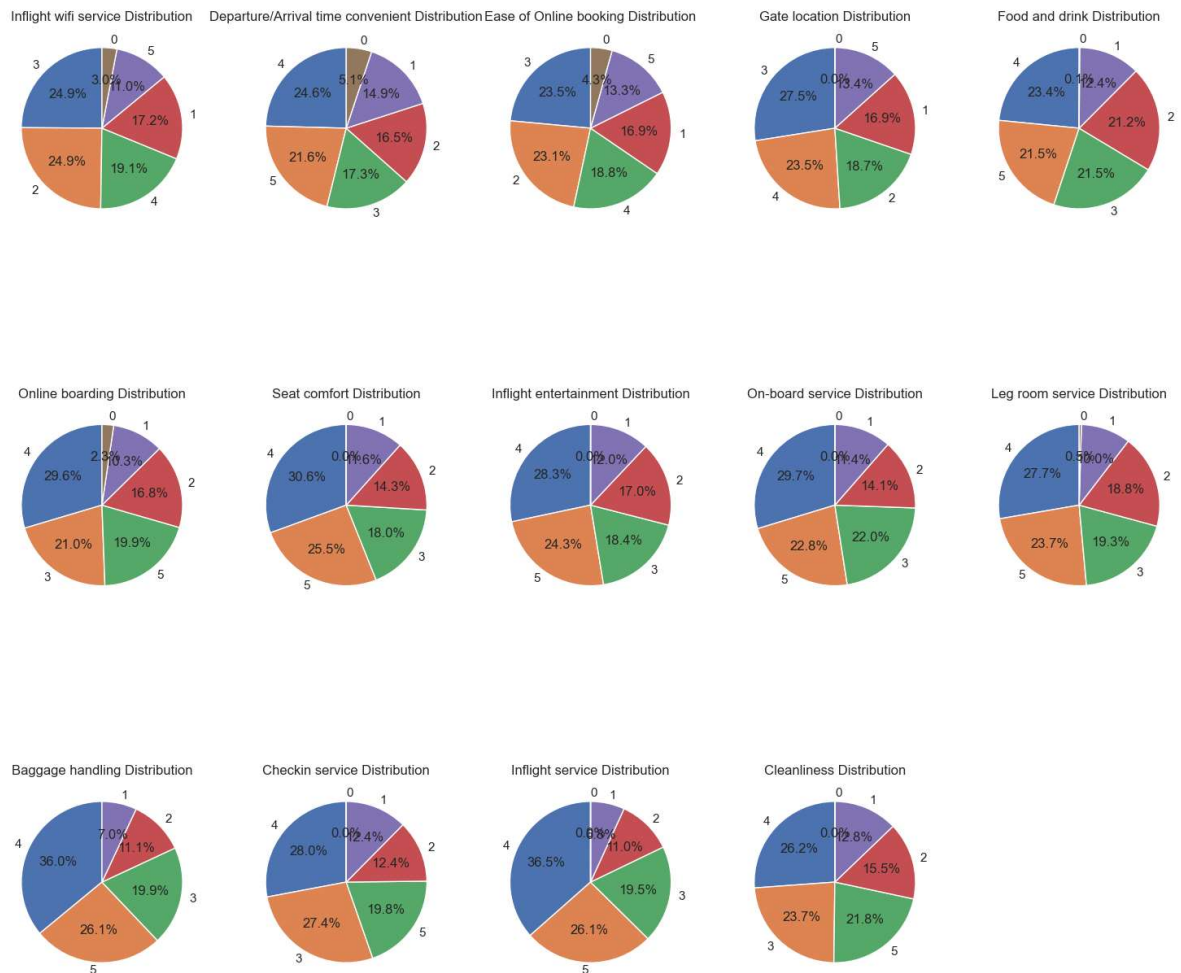
# create a pie chart for each categorical variable
for i, var in enumerate(cat_vars):
    if i < len(axs.flat):
        # count the number of occurrences for each category
        cat_counts = df[var].value_counts()

        # create a pie chart
        axs.flat[i].pie(cat_counts, labels=cat_counts.index, autopct='%1.1f%%')

        # set a title for each subplot
        axs.flat[i].set_title(f'{var} Distribution')

# adjust spacing between subplots
fig.tight_layout()
fig.delaxes(axs[2][4])
# show the plot
plt.show()

```



```

In [7]: num_vars = ['Age', 'Flight Distance', 'Departure Delay in Minutes', 'Arrival Delay in Minutes']

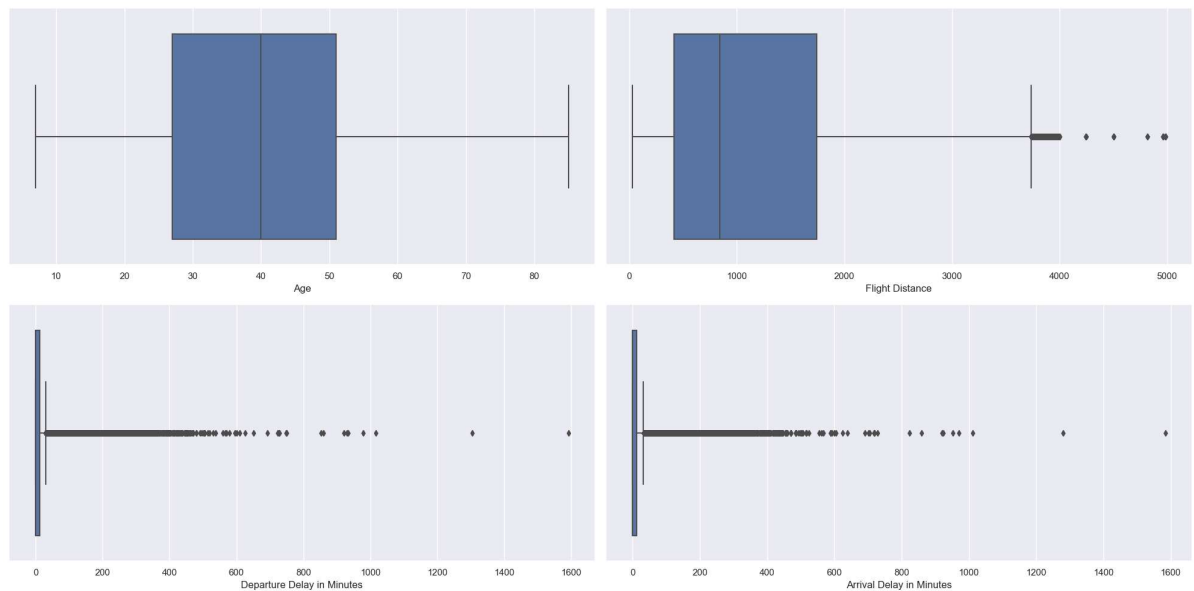
fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(20, 10))
axs = axs.flatten()

for i, var in enumerate(num_vars):
    sns.boxplot(x=var, data=df, ax=axs[i])

fig.tight_layout()

plt.show()

```

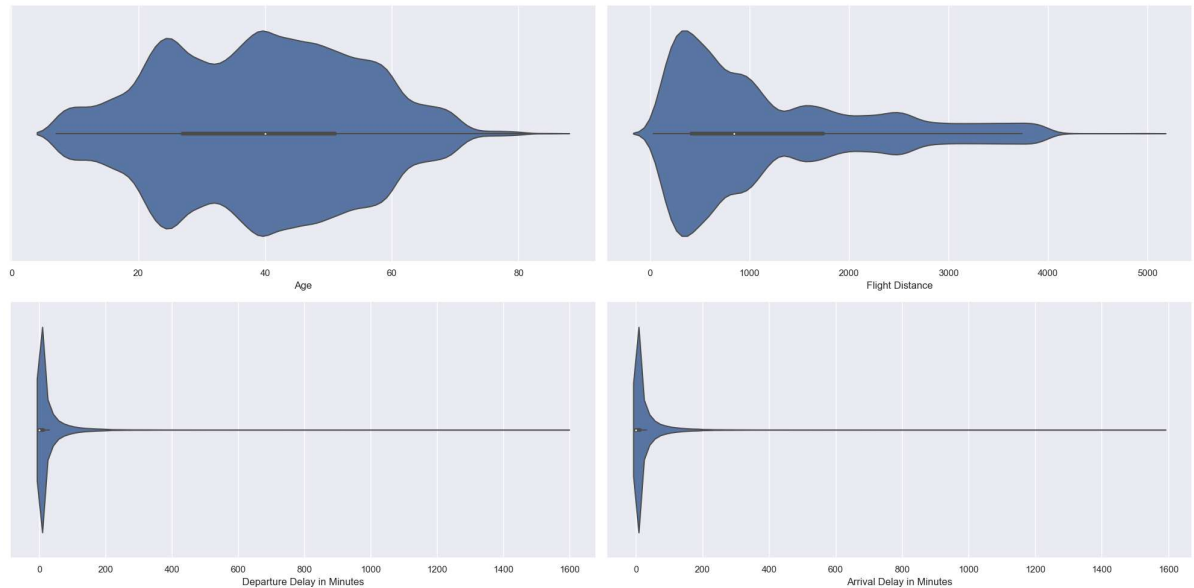


```
In [8]: num_vars = ['Age', 'Flight Distance', 'Departure Delay in Minutes', 'Arrival Delay in Minutes']

fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(20, 10))
axs = axs.flatten()

for i, var in enumerate(num_vars):
    sns.violinplot(x=var, data=df, ax=axs[i])

fig.tight_layout()
plt.show()
```



```

In [9]: num_vars = ['Age', 'Flight Distance', 'Departure Delay in Minutes', 'Arrival Delay in Minutes']

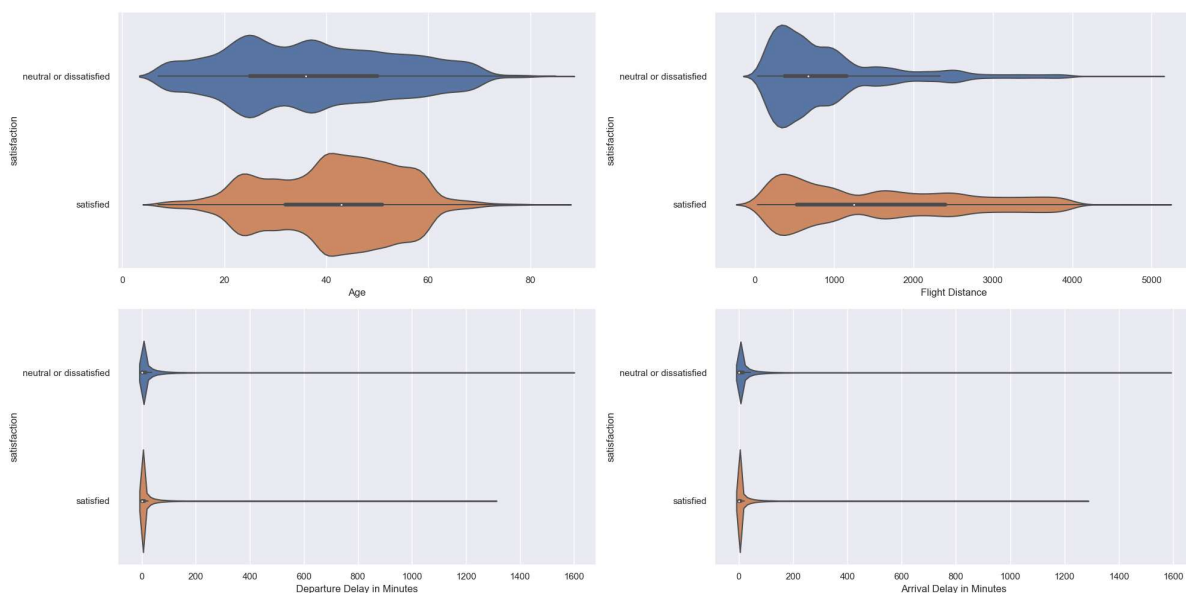
fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(20, 10))
axs = axs.flatten()

for i, var in enumerate(num_vars):
    sns.violinplot(x=var, y='satisfaction', data=df, ax=axs[i])

fig.tight_layout()

plt.show()

```



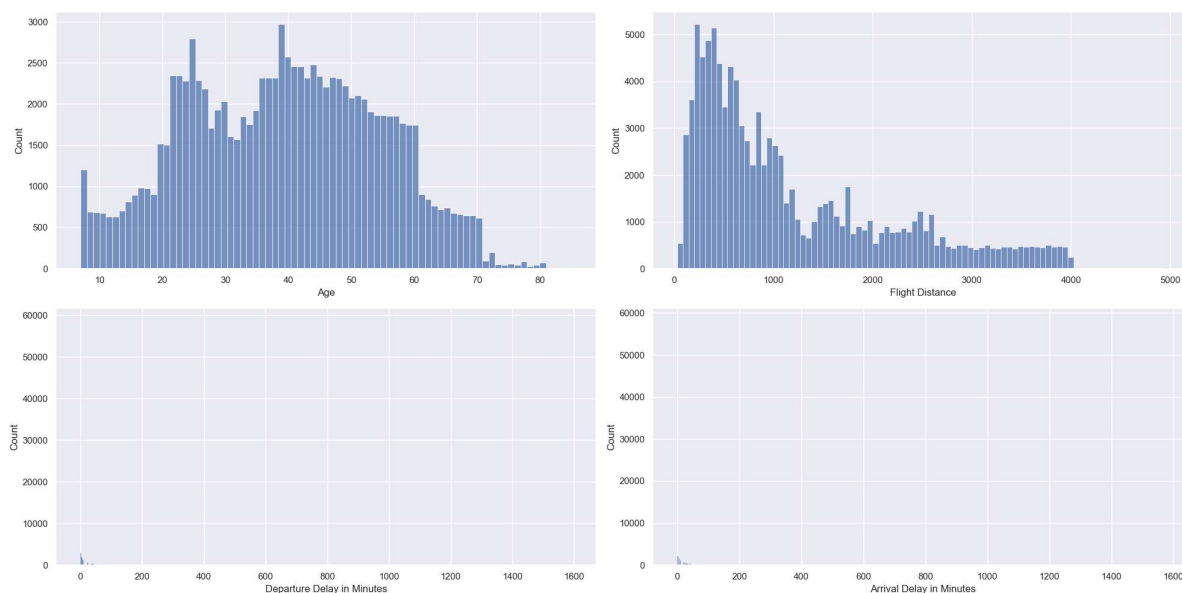

```
In [10]: num_vars = ['Age', 'Flight Distance', 'Departure Delay in Minutes', 'Arrival Delay in Minutes']

fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(20, 10))
axs = axs.flatten()


for i, var in enumerate(num_vars):
    sns.histplot(x=var, data=df, ax=axs[i])

fig.tight_layout()

plt.show()
```



```
In [11]: plt.figure(figsize=(8,8),dpi=100)
sns.scatterplot(x="Departure Delay in Minutes", y="Arrival Delay in Minutes",
plt.show())
```



Data Preprocessing

Check null value

```
In [12]: check_missing = df.isnull().sum() * 100 / df.shape[0]
check_missing[check_missing > 0].sort_values(ascending=False)
```

```
Out[12]: Arrival Delay in Minutes    0.298352
dtype: float64
```

```
In [13]: df['Arrival Delay in Minutes'].fillna(df['Arrival Delay in Minutes'].median(),
```

```
In [14]: df.dtypes
```

```
Out[14]: Gender                object
Customer Type                object
Age                          int64
Type of Travel              object
Class                      object
Flight Distance             int64
Inflight wifi service       int64
Departure/Arrival time convenient int64
Ease of Online booking      int64
Gate location              int64
Food and drink              int64
Online boarding             int64
Seat comfort                int64
Inflight entertainment      int64
On-board service            int64
Leg room service            int64
Baggage handling            int64
Checkin service             int64
Inflight service            int64
Cleanliness                 int64
Departure Delay in Minutes  int64
Arrival Delay in Minutes   float64
satisfaction                object
dtype: object
```

Show Unique Value each Columns

```
In [15]: # Loop over each column in the DataFrame where dtype is 'object'
for col in df.select_dtypes(include=['object']).columns:

    # Print the column name and the unique values
    print(f"{col}: {df[col].unique()}")
```

```
Gender: ['Male' 'Female']
Customer Type: ['Loyal Customer' 'disloyal Customer']
Type of Travel: ['Personal Travel' 'Business travel']
Class: ['Eco Plus' 'Business' 'Eco']
satisfaction: ['neutral or dissatisfied' 'satisfied']
```

Label encoding each categorical column

```
In [16]: from sklearn import preprocessing

# Loop over each column in the DataFrame where dtype is 'object'
for col in df.select_dtypes(include=['object']).columns:

    # Initialize a LabelEncoder object
    label_encoder = preprocessing.LabelEncoder()

    # Fit the encoder to the unique values in the column
    label_encoder.fit(df[col].unique())

    # Transform the column using the encoder
    df[col] = label_encoder.transform(df[col])


    # Print the column name and the unique encoded values
    print(f"{col}: {df[col].unique()}")
```

```
Gender: [1 0]
Customer Type: [0 1]
Type of Travel: [1 0]
Class: [2 0 1]
satisfaction: [0 1]
```

Check class value

```
In [17]: sns.countplot(df['satisfaction'])  
df['satisfaction'].value_counts()
```


```
Out[17]: 0    58879  
        1    45025  
        Name: satisfaction, dtype: int64
```



```
In [18]: from sklearn.utils import resample  
#create two different dataframe of majority and minority class  
df_majority = df[(df['satisfaction']==0)]  
df_minority = df[(df['satisfaction']==1)]  
# upsample minority class  
df_minority_upsampled = resample(df_minority,  
                                replace=True,      # sample with replacement  
                                n_samples= 58879, # to match majority class  
                                random_state=0)   # reproducible results  
# Combine majority class with upsampled minority class  
df_upsampled = pd.concat([df_minority_upsampled, df_majority])
```

```
In [19]: sns.countplot(df_upsampled['satisfaction'])  
df_upsampled['satisfaction'].value_counts()
```

```
Out[19]: 1    58879  
0    58879  
Name: satisfaction, dtype: int64
```



Remove Outlier using IQR

```

In [20]: # specify the columns to remove outliers from dataframe
column_names = ['Age', 'Flight Distance', 'Departure Delay in Minutes', 'Arrival Delay in Minutes']

# remove outliers for each selected column using the IQR method
for column_name in column_names:
    Q1 = df_upsampled[column_name].quantile(0.25)
    Q3 = df_upsampled[column_name].quantile(0.75)
    IQR = Q3 - Q1
    df_upsampled = df_upsampled[~((df_upsampled[column_name] < (Q1 - 1.5 * IQR) | df_upsampled[column_name] > (Q3 + 1.5 * IQR)))]

df_upsampled.head()

```

Out[20]:


	Gender	Customer Type	Age	Type of Travel	Class	Flight Distance	Inflight wifi service	Departure/Arrival time convenient	Ease of Online booking	log
6287	0	0	43	0	0	3603	3	3	3	
100566	1	0	45	0	1	451	4	2	2	
98330	1	0	31	0	0	1334	1	1	1	
48752	0	0	51	0	1	589	4	2	2	
69983	1	0	46	0	0	1400	2	2	5	

```
In [21]: df_upsampled.shape
```

Out[21]: (87448, 23)

```
In [22]: plt.figure(figsize=(15,12))
sns.heatmap(df_upsampled.corr(), fmt='.2g')
```

Out[22]: <AxesSubplot:>



Machine Learning Model Building

```
In [23]: X = df_upsampled.drop('satisfaction', axis=1)
y = df_upsampled['satisfaction']
```

```
In [24]: #test size 20% and train size 80%
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,random_
```

Decision Tree


```
In [25]: from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import GridSearchCV
dtree = DecisionTreeClassifier()
param_grid = {
    'max_depth': [3, 4, 5, 6, 7, 8],
    'min_samples_split': [2, 3, 4],
    'min_samples_leaf': [1, 2, 3, 4]
}

# Perform a grid search with cross-validation to find the best hyperparameters
grid_search = GridSearchCV(dtree, param_grid, cv=5)
grid_search.fit(X_train, y_train)

# Print the best hyperparameters
print(grid_search.best_params_)

{'max_depth': 8, 'min_samples_leaf': 4, 'min_samples_split': 2}
```

```
In [26]: from sklearn.tree import DecisionTreeClassifier
dtree = DecisionTreeClassifier(random_state=0, max_depth=8, min_samples_leaf=4)
dtree.fit(X_train, y_train)
```

Out[26]: DecisionTreeClassifier(max_depth=8, min_samples_leaf=4, random_state=0)

```
In [27]: y_pred = dtree.predict(X_test)
print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100, 2), "%")
```


Accuracy Score : 93.62 %

```
In [28]: from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
print('F-1 Score : ', (f1_score(y_test, y_pred, average='micro')))
print('Precision Score : ', (precision_score(y_test, y_pred, average='micro')))
print('Recall Score : ', (recall_score(y_test, y_pred, average='micro')))
print('Jaccard Score : ', (jaccard_score(y_test, y_pred, average='micro')))
print('Log Loss : ', (log_loss(y_test, y_pred)))
```

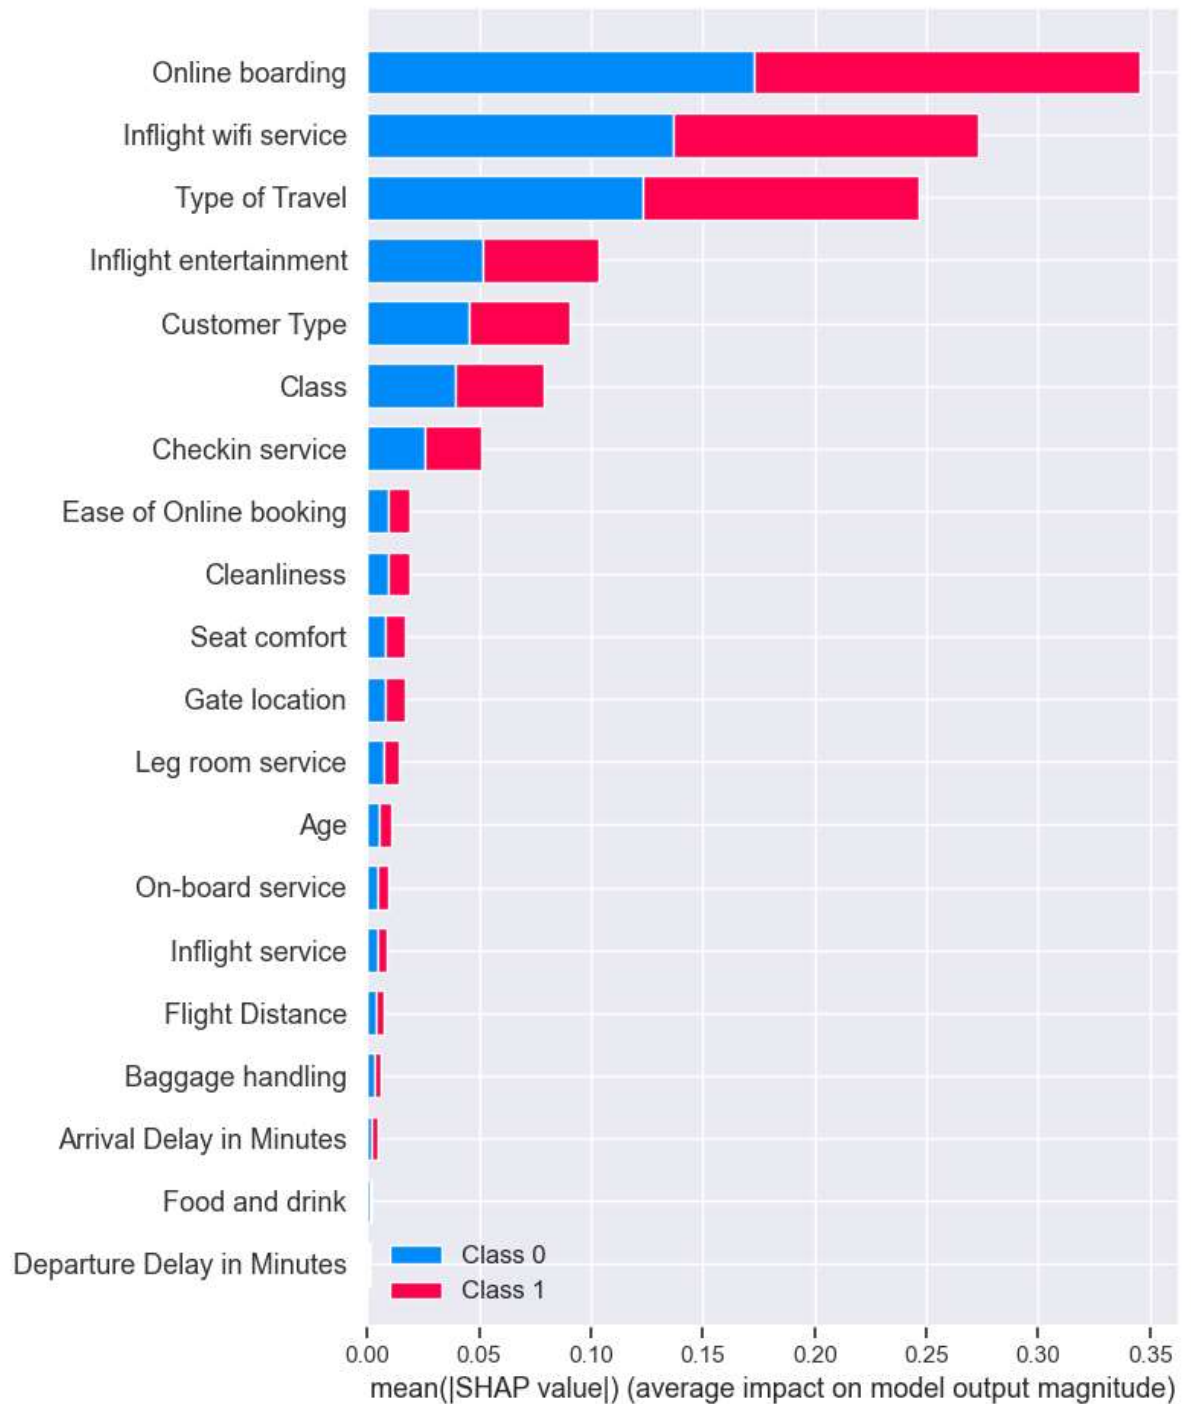
F-1 Score : 0.9361921097770154
Precision Score : 0.9361921097770154
Recall Score : 0.9361921097770154
Jaccard Score : 0.8800386971944534
Log Loss : 2.2038680769756898

```
In [29]: imp_df = pd.DataFrame({
    "Feature Name": X_train.columns,
    "Importance": dtree.feature_importances_
})
fi = imp_df.sort_values(by="Importance", ascending=False)


fi2 = fi.head(10)
plt.figure(figsize=(10,8))
sns.barplot(data=fi2, x='Importance', y='Feature Name')
plt.title('Top 10 Feature Importance Each Attributes (Decision Tree)', fontsize=14)
plt.xlabel ('Importance', fontsize=16)
plt.ylabel ('Feature Name', fontsize=16)
plt.show()
```



```
In [30]: import shap
explainer = shap.TreeExplainer(dtree)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values, X_test)
```



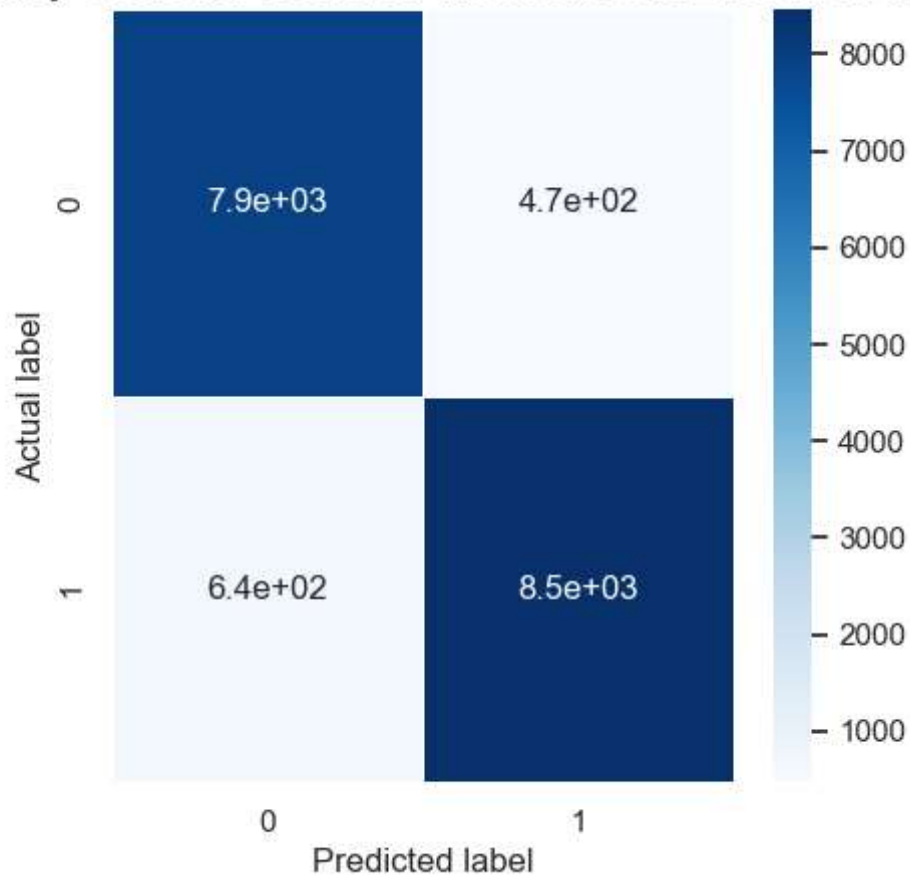
```
In [31]: # compute SHAP values
explainer = shap.TreeExplainer(dtree)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values[1], X_test.values, feature_names = X_test.columns)
```



```
In [32]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5,5))
sns.heatmap(data=cm,linewidths=.5, annot=True, cmap = 'Blues')
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
all_sample_title = 'Accuracy Score for Decision Tree: {0}'.format(dtree.score())
plt.title(all_sample_title, size = 15)
```

Out[32]: Text(0.5, 1.0, 'Accuracy Score for Decision Tree: 0.9361921097770154')

Accuracy Score for Decision Tree: 0.9361921097770154




```
In [33]: from sklearn.metrics import roc_curve, roc_auc_score
y_pred_proba = dtree.predict_proba(X_test)[:][:,1]

df_actual_predicted = pd.concat([pd.DataFrame(np.array(y_test), columns=['y_actual'], index=y_test.index),
                                pd.DataFrame(y_pred_proba, columns=['y_predicted'], index=y_test.index)])

fpr, tpr, tr = roc_curve(df_actual_predicted['y_actual'], df_actual_predicted['y_predicted'])
auc = roc_auc_score(df_actual_predicted['y_actual'], df_actual_predicted['y_predicted'])

plt.plot(fpr, tpr, label='AUC = %0.4f' %auc)
plt.plot(fpr, fpr, linestyle='--', color='k')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve', size = 15)
plt.legend()
```

Out[33]: <matplotlib.legend.Legend at 0x229b4b3d8e0>



Random Forest

```
In [34]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
rfc = RandomForestClassifier()
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [None, 5, 10],
    'max_features': ['sqrt', 'log2', None]
}

# Perform a grid search with cross-validation to find the best hyperparameters
grid_search = GridSearchCV(rfc, param_grid, cv=5)
grid_search.fit(X_train, y_train)

# Print the best hyperparameters
print(grid_search.best_params_)

{'max_depth': None, 'max_features': 'sqrt', 'n_estimators': 200}
```

```
In [35]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier(random_state=0, max_features='sqrt', n_estimators=200)
rfc.fit(X_train, y_train)
```

Out[35]: RandomForestClassifier(max_features='sqrt', n_estimators=200, random_state=0)

```
In [36]: y_pred = rfc.predict(X_test)
print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")

Accuracy Score : 97.61 %
```

```
In [37]: from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
print('F-1 Score : ',(f1_score(y_test, y_pred, average='micro')))
print('Precision Score : ',(precision_score(y_test, y_pred, average='micro')))
print('Recall Score : ',(recall_score(y_test, y_pred, average='micro')))
print('Jaccard Score : ',(jaccard_score(y_test, y_pred, average='micro')))
print('Log Loss : ',(log_loss(y_test, y_pred)))

F-1 Score : 0.9761006289308176
Precision Score : 0.9761006289308176
Recall Score : 0.9761006289308176
Jaccard Score : 0.9533169533169533
Log Loss : 0.8254636282098088
```

```
In [38]: imp_df = pd.DataFrame({
    "Feature Name": X_train.columns,
    "Importance": rfc.feature_importances_
})
fi = imp_df.sort_values(by="Importance", ascending=False)

fi2 = fi.head(10)
plt.figure(figsize=(10,8))
sns.barplot(data=fi2, x='Importance', y='Feature Name')
plt.title('Top 10 Feature Importance Each Attributes (Random Forest)', fontsize=14)
plt.xlabel ('Importance', fontsize=16)
plt.ylabel ('Feature Name', fontsize=16)
plt.show()
```

